

# Chapel: Productive Parallel Programming at Scale

Brad Chamberlain, Chapel Team, Cray Inc.

Northwest C++ Users Group

October 19<sup>th</sup>, 2016



# Safe Harbor Statement

This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.



# Cray: The Supercomputing Company

**1972: Seymour Cray founded Cray Research**

**2000: Tera purchased Cray Research from SGI and formed Cray Inc.**

- corporate headquarters based in Seattle, WA

## Technology Focus Areas:

- Computation
- Storage
- Analytics

**Vision: Provide the systems and tools that our customers need to solve the world's hardest problems.**



# High Performance Computing (HPC) Programming Models by Example



COMPUTE

| STORE

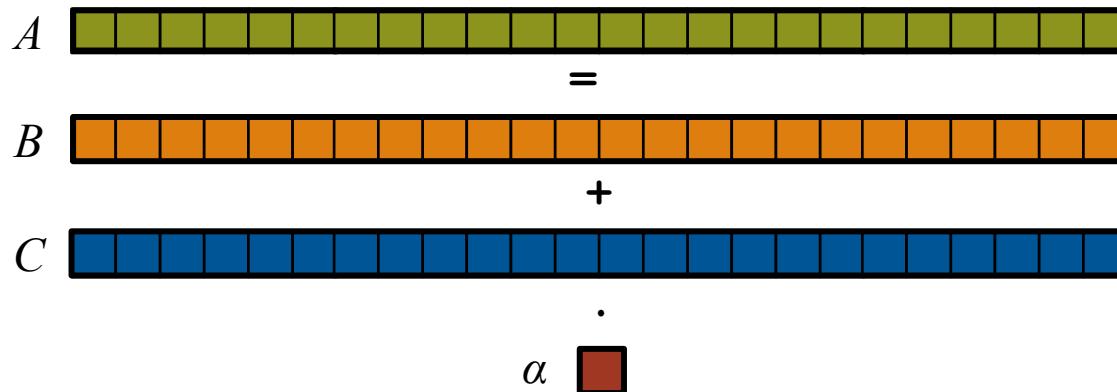
| ANALYZE

# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures:**

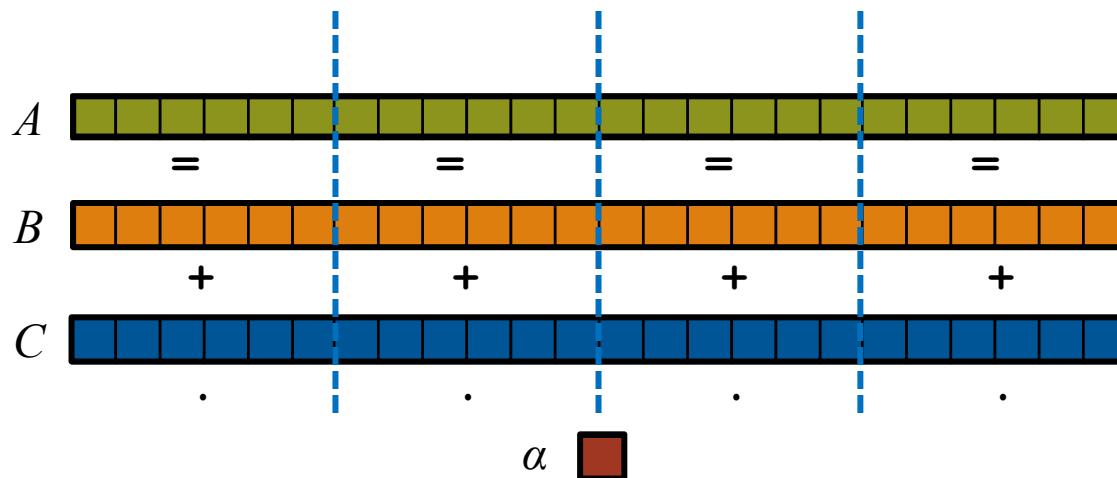


# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures, in parallel:**

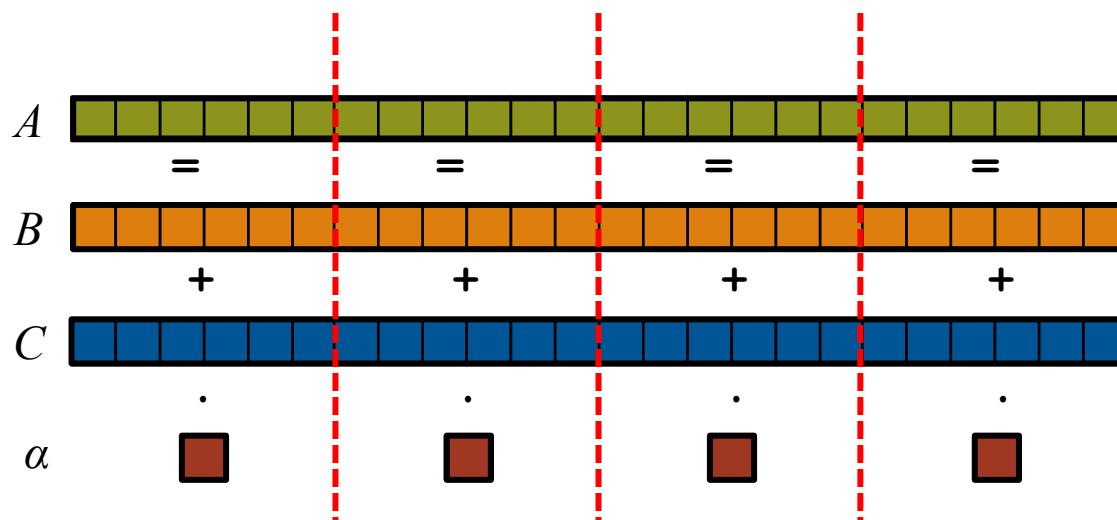


# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures, in parallel (distributed memory):**

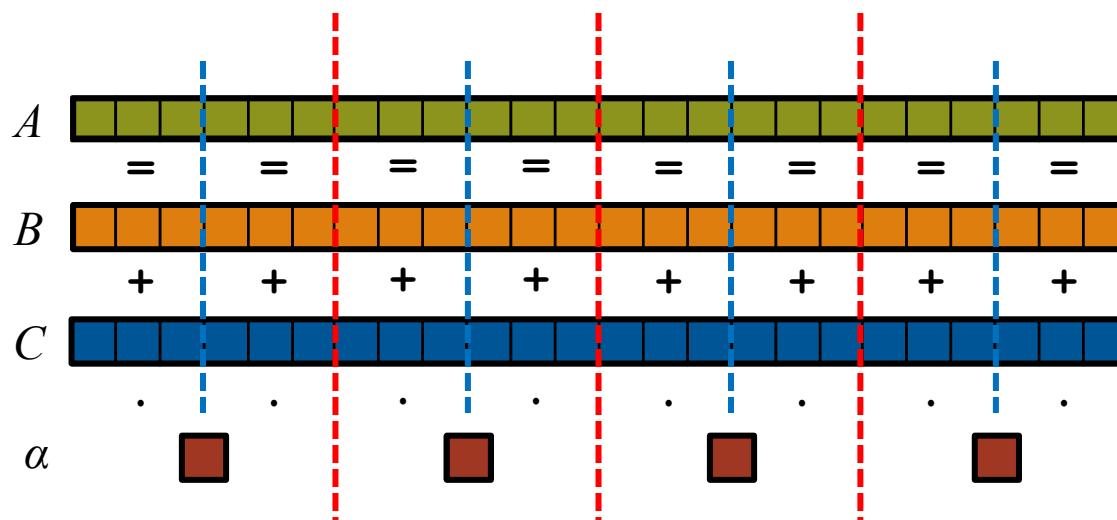


# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures, in parallel (distributed memory multicore):**



# STREAM Triad: MPI

```
#include <hpcc.h>

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
                0, comm );
}

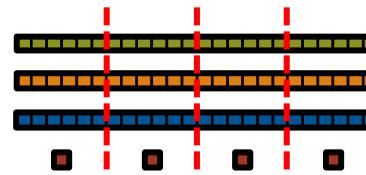
return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
                                       sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
}
```

MPI



```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory
(%d).\n", VectorSize );
        fclose( outFile );
    }
    return 1;
}

for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 0.0;
}

scalar = 3.0;

for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);
```



# STREAM Triad: MPI+OpenMP

## MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
                0, comm );

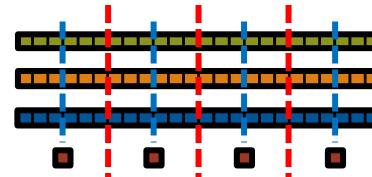
    return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
                                       sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
}

```



```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory
(%d).\n", VectorSize );
        fclose( outFile );
    }
    return 1;
}

#ifndef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 0.0;
}

scalar = 3.0;

#ifndef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);

```

# STREAM Triad: MPI+OpenMP vs. CUDA

## MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );
    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );
    return errCount;
}

int HPCC_LocalVectorSize(HPCC_Params *params, int len)
{
    double scalar;
    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );
    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }

    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }
    scalar = 3.0;

    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];

    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
    return 0;
}
```

## CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    if( N % dimBlock.x != 0 ) dimGrid

    set_array<<<dimGrid,dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid,dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid,dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

    __global__ void set_array(float *a, float value, int len) {
        int idx = threadIdx.x + blockIdx.x * blockDim.x;
        if (idx < len) a[idx] = value;
    }

    __global__ void STREAM_Triad( float *a, float *b, float *c,
                                float scalar, int len) {
        int idx = threadIdx.x + blockIdx.x * blockDim.x;
        if (idx < len) c[idx] = a[idx]+scalar*b[idx];
    }
}
```

COMPUTE

STORE

ANALYZE

# Why so many programming models?

HPC tends to approach programming models bottom-up:

Given a system and its core capabilities...

...provide features that can access the available performance.

- portability? generality? programmability? ...not strictly required.

Type of HW Parallelism	Programming Model	Unit of Parallelism
Inter-node	MPI	executable
Intra-node/multicore	OpenMP / pthreads	iteration/task
Instruction-level vectors/threads	pragmas	iteration
GPU/accelerator	CUDA / Open[CL MP ACC]	SIMD function/task

**benefits:** lots of control; decent generality; easy to implement  
**downsides:** lots of user-managed detail; brittle to changes

# Motivation for Chapel

**Q: Can a single language be...**

- ...as productive as Python?
- ...as fast as Fortran?
- ...as portable as C?
- ...as scalable as MPI?
- ...as fun as <your favorite language here>?

**A: We believe so.**



# The Challenge

**Q: So why don't we have such languages already?**

**A: Technical challenges?**

- while they exist, we don't think this is the main issue...

**A: Due to a lack of...**

- ...long-term efforts
- ...resources
- ...community will
- ...co-design between developers and users
- ...patience

***Chapel is our attempt to reverse this trend***

# What is Chapel?

**Chapel:** A productive parallel programming language

- extensible
- portable
- open-source
- a collaborative effort
- a work-in-progress

## Goals:

- Support general parallel programming
  - “any parallel algorithm on any parallel hardware”
- Make parallel programming far more productive



# What does “Productivity” mean to you?

## Recent Graduates:

“something similar to what I used in school: Python, Matlab, Java, ...”

## Seasoned HPC Programmers:

“that sugary stuff that I don’t need because I ~~was born to suffer~~  
want full control  
to ensure performance”

## Computational Scientists:

“something that lets me express my parallel computations  
without having to wrestle with architecture-specific details”

## Chapel Team:

“something that lets computational scientists express what they want,  
without taking away the control that HPC programmers need,  
implemented in a language as attractive as recent graduates want.”



# Rewinding a few slides...

## MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );
    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );
    return errCount;
}

int HPCC_Triad(HPCC_Params *params, FILE *outFile)
{
    int i, j, k;
    double scalar;
    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );
    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }
    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }
    scalar = 3.0;
    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];
    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
    return 0;
}
```

## CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    if( N % dimBlock.x != 0 ) dimGrid

    set_array<<<dimGrid,dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid,dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid,dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}

__global__ void STREAM_Triad( float *a, float *b, float *c,
                             float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
}
```

COMPUTE

STORE

ANALYZE

# STREAM Triad: Chapel

MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params,
int myRank, commSize;
int rv, errCount;
MPI_Comm comm = MPI_COMM_WORLD;
MPI_Comm size( comm, &commSize );
MPI_Comm_rank( comm, &myRank );
rv = HPCC_Stream( params, 0 == myRank );
MPI_Reduce( &rv, &errCount, 1, MPI_
return errCount;
}

int HPCC_Stream(HPCC_Params *params,
register int j;
double scalar;
VectorSize = HPCC_LocalVectorSize();
a = HPCC_XMALLOC( double, VectorSize );
b = HPCC_XMALLOC( double, VectorSize );
c = HPCC_XMALLOC( double, VectorSize );

if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d).\n" VectorSize );
        fclose( outFile );
    }
}
```

Chapel

```
config const m = 1000,
alpha = 3.0;

const ProblemSpace = {1..m} dmapped ...;

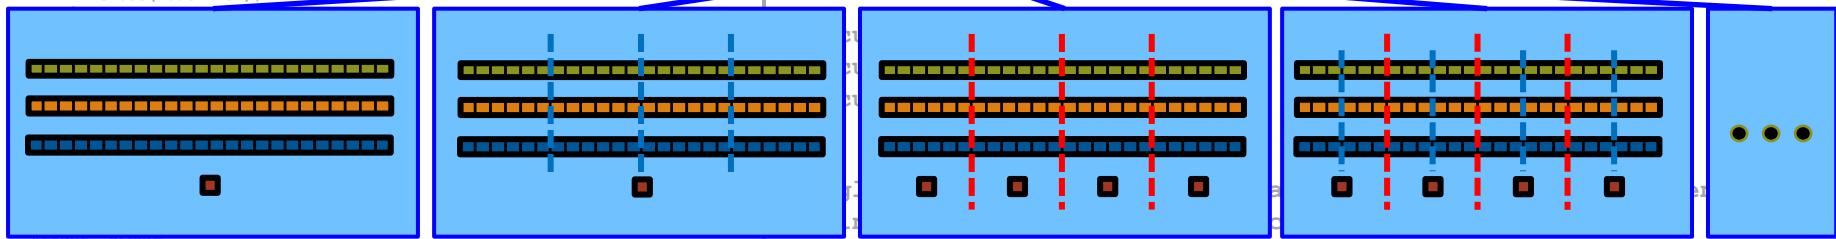
var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

A = B + alpha * C;
```

the special sauce

```
; N);
N);
l_c, d_a, scalar, N);
```



Philosophy: Good, *top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

# Outline

- ✓ Motivation for Chapel
- Survey of Chapel Concepts
  - Chapel Project and Characterizations
  - Chapel Resources

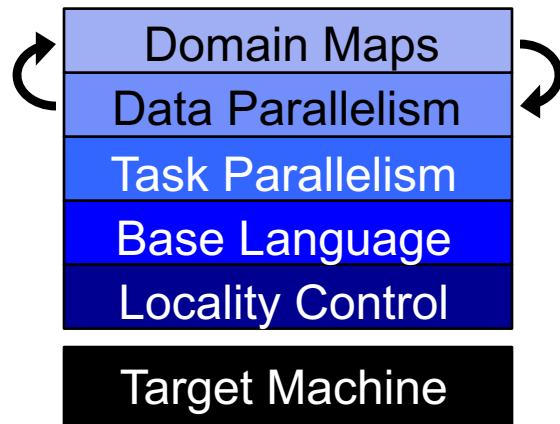


# Chapel's Multiresolution Philosophy

## ***Multiresolution Design:*** Support multiple tiers of features

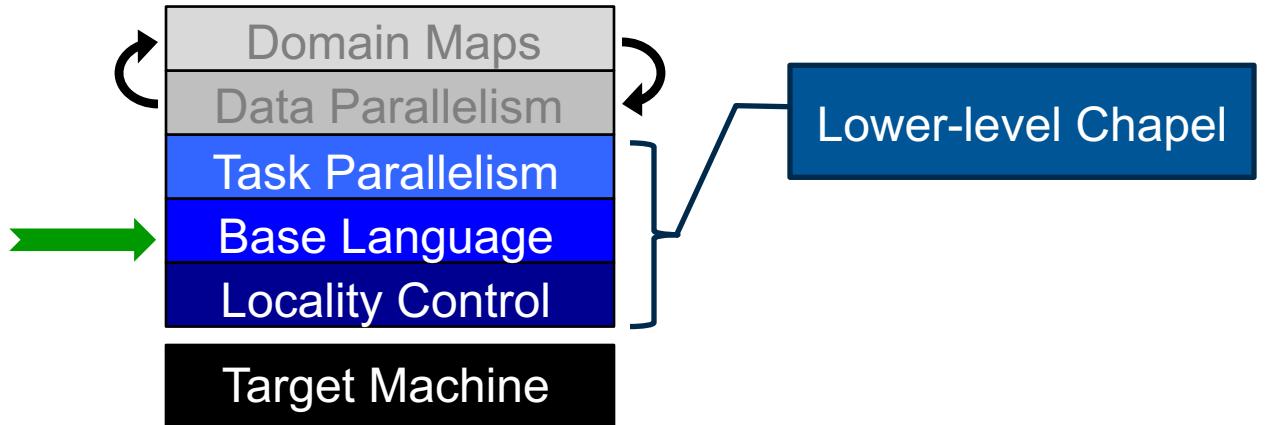
- higher levels for programmability, productivity
- lower levels for greater degrees of control

*Chapel language concepts*



- build the higher-level concepts in terms of the lower
- permit the user to intermix layers arbitrarily

# Base Language



# Base Language Features, by example

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for f in fib(n) do
    writeln(f);
```

```
0  
1  
1  
2  
3  
5  
8  
...
```

# Base Language Features, by example

## CLU-style iterators

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=gt; next;
    }
}
```

```
config const n = 10;

for f in fib(n) do
    writeln(f);
```

```
0  
1  
1  
2  
3  
5  
8  
...
```

# Base Language Features, by example

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

Configuration declarations  
(to avoid command-line argument parsing)  
./a.out -n=1000000

```
config const n = 10;

for f in fib(n) do
    writeln(f);
```

```
0
1
1
2
3
5
8
...
...
```

# Base Language Features, by example

Static type inference for:

- arguments
- return types
- variables

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

for f in fib(n) do
    writeln(f);
```

```
0
1
1
2
3
5
8
...
```

# Base Language Features, by example

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

**config const n = 10;**

```
for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

**fib #0 is 0**  
**fib #1 is 1**  
**fib #2 is 1**  
**fib #3 is 2**  
**fib #4 is 3**  
**fib #5 is 5**  
**fib #6 is 8**  
...

Zippered iteration

# Base Language Features, by example

## Range types and operators

```
iter fib(n) {
    var current = 0
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=gt; next;
    }
}
```

```
config const n = 10;

for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```

# Base Language Features, by example

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

tuples

```
config const n = 10;

for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```

# Base Language Features, by example

```
iter fib(n) {
    var current = 0,
        next = 1;

    for i in 1..n {
        yield current;
        current += next;
        current <=> next;
    }
}
```

```
config const n = 10;

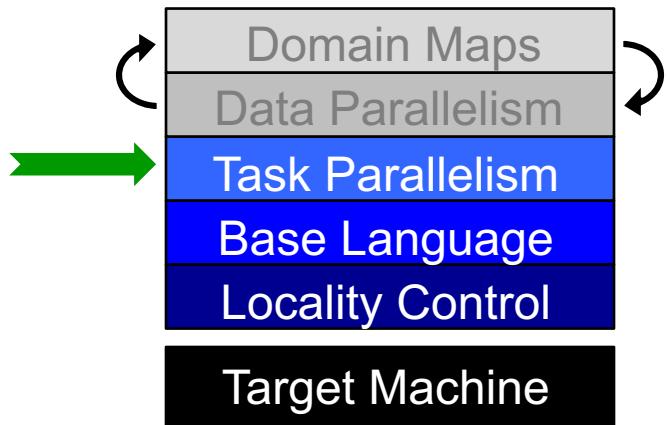
for (i,f) in zip(0..#n, fib(n)) do
    writeln("fib #", i, " is ", f);
```

```
fib #0 is 0
fib #1 is 1
fib #2 is 1
fib #3 is 2
fib #4 is 3
fib #5 is 5
fib #6 is 8
...
```

# Other Base Language Features

- **interoperability features**
- **OOP** (value- and reference-based)
- **overloading, where clauses**
- **argument intents, default values, match-by-name**
- **compile-time features for meta-programming**
  - e.g., compile-time functions to compute types, values; reflection
- **modules** (for namespace management)
- **rank-independent programming features**
- ...

# Task Parallelism



# Task Parallelism: Begin Statements

```
// create a fire-and-forget task for a statement
begin writeln("hello world");
writeln("goodbye");
```

## Possible outputs:

hello world  
goodbye

goodbye  
hello world

# Task Parallelism: Coforall Loops

```
// create a task per iteration
coforall t in 0..#numTasks {
    writeln("Hello from task ", t, " of ", numTasks);
} // implicit join of the numTasks tasks here

writeln("All tasks done");
```

## Sample output:

```
Hello from task 2 of 4
Hello from task 0 of 4
Hello from task 3 of 4
Hello from task 1 of 4
All tasks done
```



# Task Parallelism: Data-Driven Synchronization

- **atomic variables:** support atomic operations
  - e.g., compare-and-swap; atomic sum, multiply, etc.
  - similar to C/C++
  
- **sync variables:** store full-empty state along with value
  - by default, reads/writes block until full/empty, leave in opposite state

# Bounded Buffer Producer/Consumer Example

```
begin producer();
consumer();

// 'sync' types store full/empty state along with value
var buff$: [0..#buffersize] sync real;

proc producer() {
    var i = 0;
    for ... {
        i = (i+1) % buffersize;
        buff$[i] = ...; // writes block until empty, leave full
    }
}

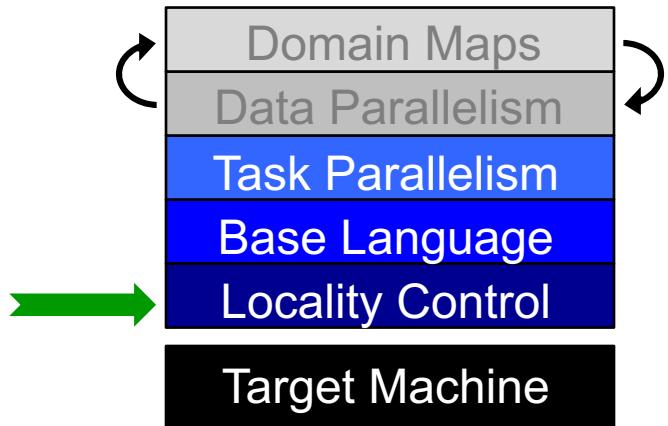
proc consumer() {
    var i = 0;
    while ... {
        i= (i+1) % buffersize;
        ...buff$[i]...; // reads block until full, leave empty
    }
}
```



# Other Task Parallel Concepts

- **cobegins**: create tasks using compound statements
- **single variables**: like sync variables, but write-once
- **sync statements**: join unstructured tasks
- **serial statements**: conditionally squash parallelism

# Locality Control



# The Locale Type

## Definition:

- Abstract unit of target architecture
- Supports reasoning about locality
  - defines “here vs. there” / “local vs. remote”
- Capable of running tasks and storing variables
  - i.e., has processors and memory

**Typically:** A compute node (multicore processor or SMP)



# Getting started with locales

- Specify # of locales when running Chapel programs

```
% a.out --numLocales=8
```

```
% a.out -nl 8
```

- Chapel provides built-in locale variables

```
config const numLocales: int = ...;  
const Locales: [0..#numLocales] locale = ...;
```

*Locales*



- main () starts execution as a task on locale #0



# Locale Operations

- Locale methods support queries about the target system:

```
proc locale.physicalMemory(...) { ... }  
proc locale.numCores { ... }  
proc locale.id { ... }  
proc locale.name { ... }
```

- On-clauses support placement of computations:

```
writeln("on locale 0");  
on Locales[1] do  
    writeln("now on locale 1");  
writeln("on locale 0 again");
```

```
on A[i,j] do  
    bigComputation(A);  
  
on node.left do  
    search(node.left);
```

# Parallelism and Locality: Orthogonal in Chapel

- This is a **parallel**, but local program:

```
coforall i in 1..msgs do  
    writeln("Hello from task ", i);
```

- This is a **distributed**, but serial program:

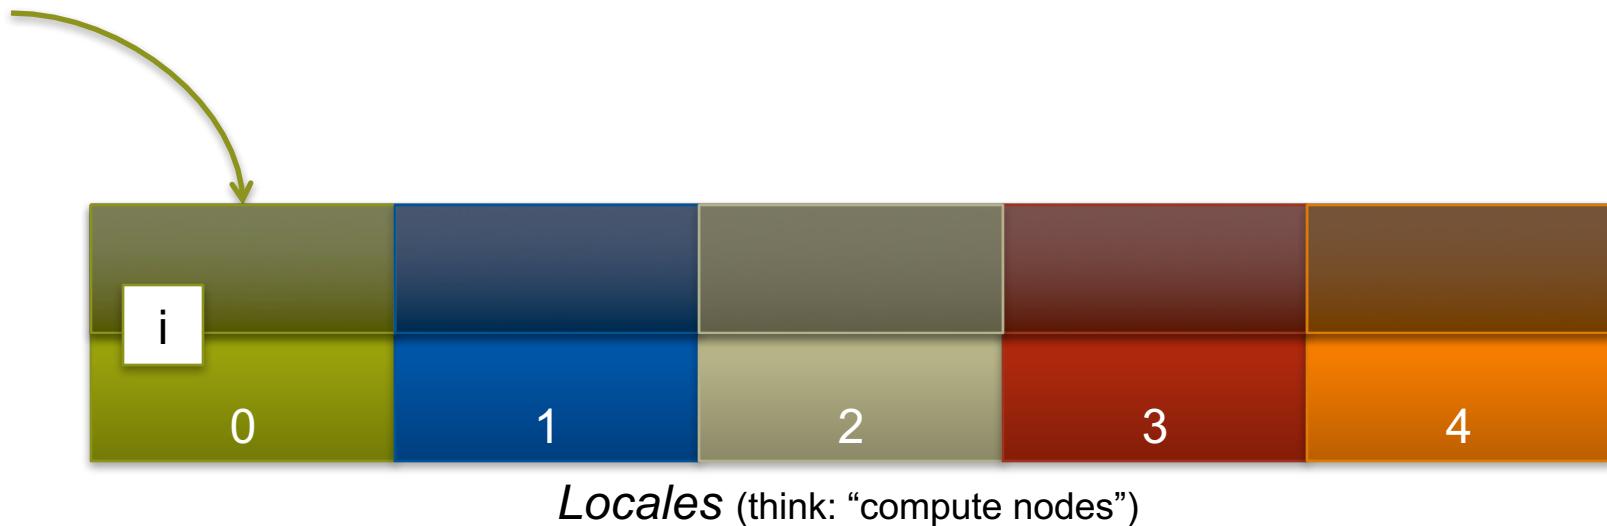
```
writeln("Hello from locale 0!");  
on Locales[1] do writeln("Hello from locale 1!");  
on Locales[2] do writeln("Hello from locale 2!");
```

- This is a **distributed parallel** program:

```
coforall i in 1..msgs do  
    on Locales[i%numLocales] do  
        writeln("Hello from task ", i,  
               " running on locale ", here.id);
```

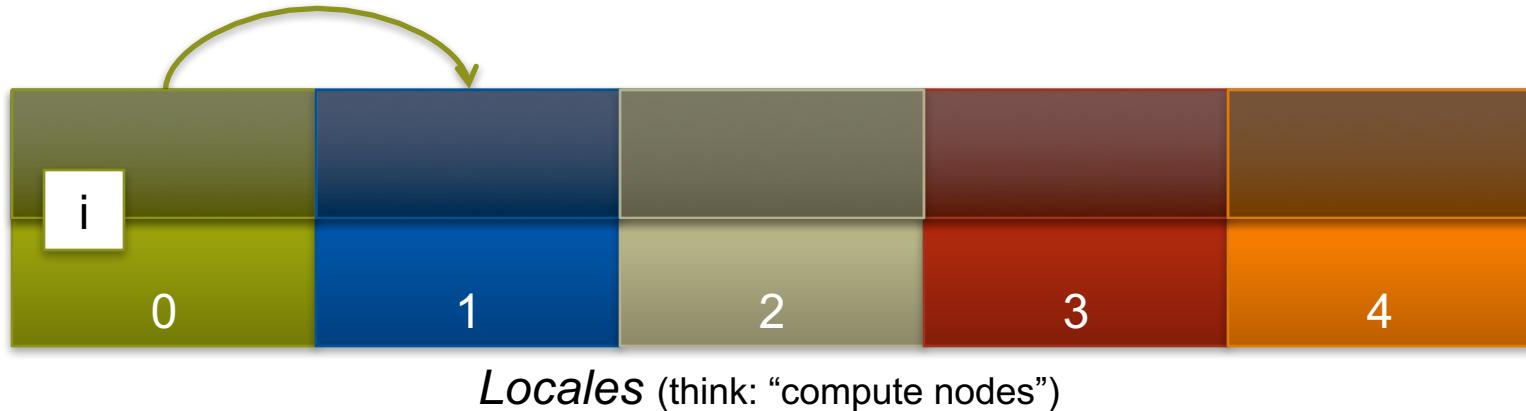
# Chapel: Scoping and Locality

```
var i: int;
```



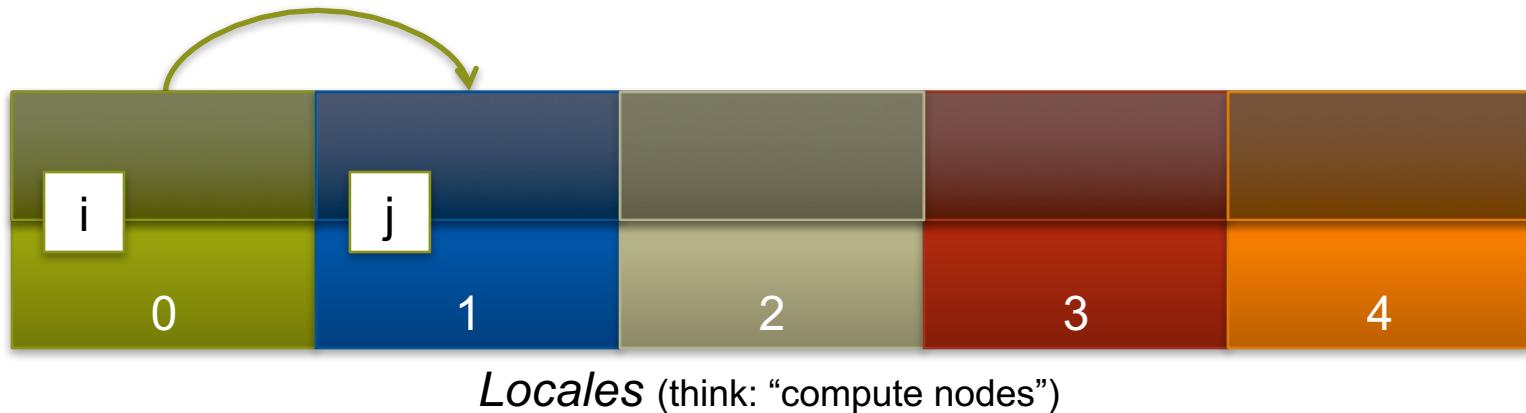
# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {
```



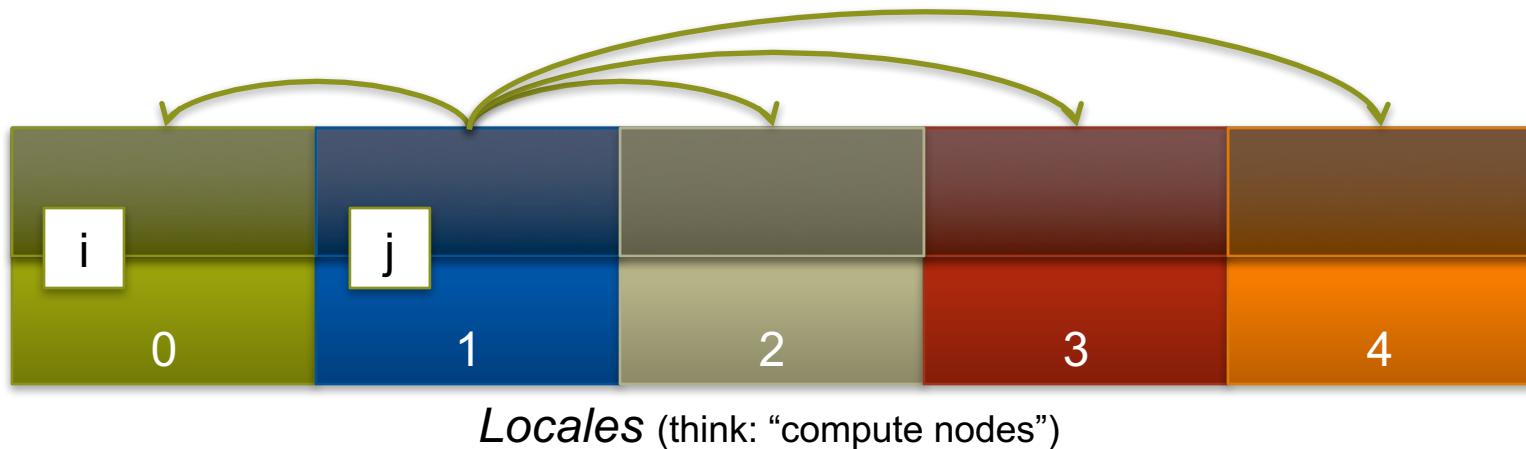
# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;
```



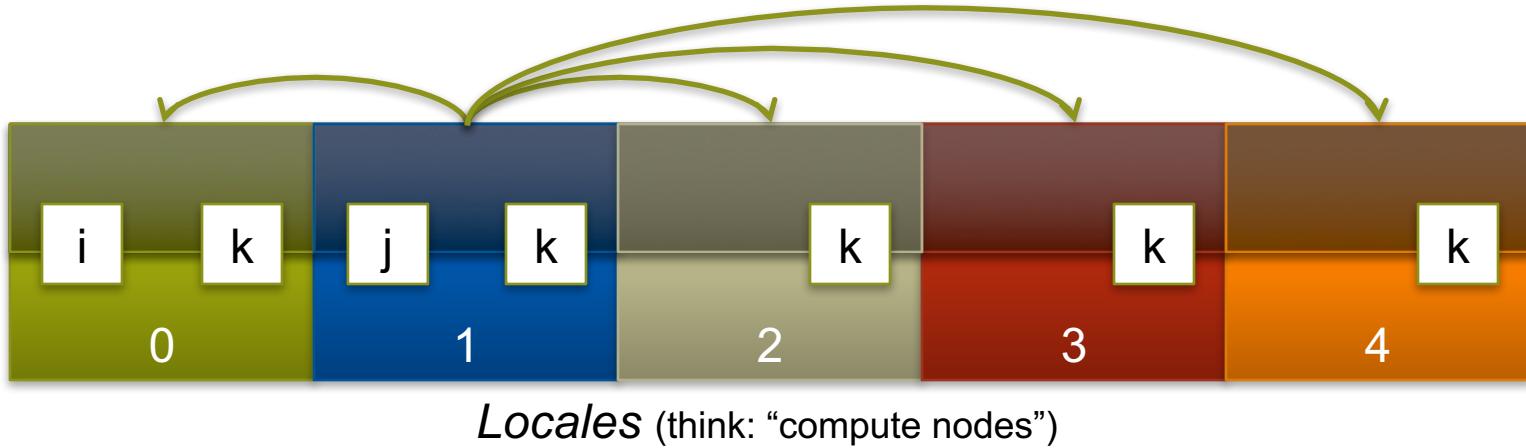
# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {
```



# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
            ...  
        }  
    }  
}
```

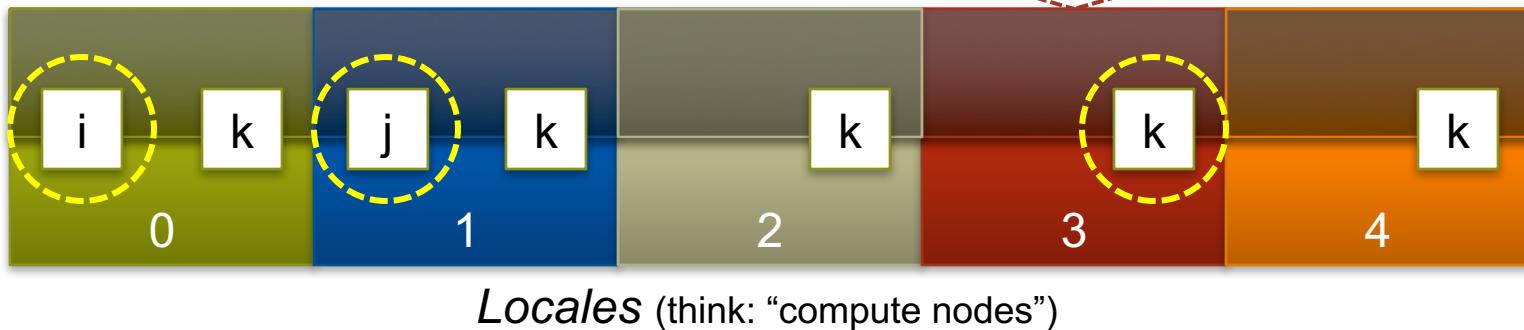


# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
            k = 2*i + j;  
        }  
    }  
}
```

OK to access  $i$ ,  $j$ , and  $k$  wherever they live

$k = 2*i + j;$

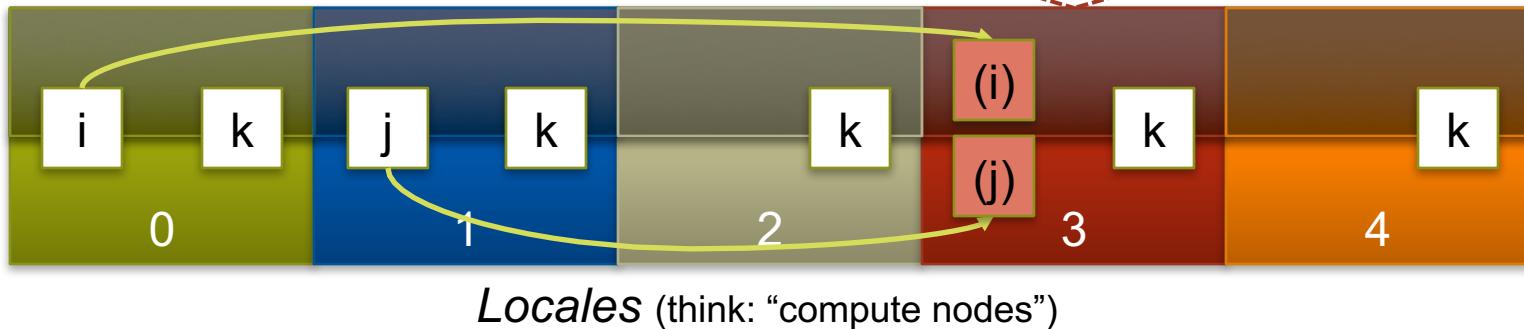


# Chapel: Scoping and Locality

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
            k = 2*i + j;  
        }  
    }  
}
```

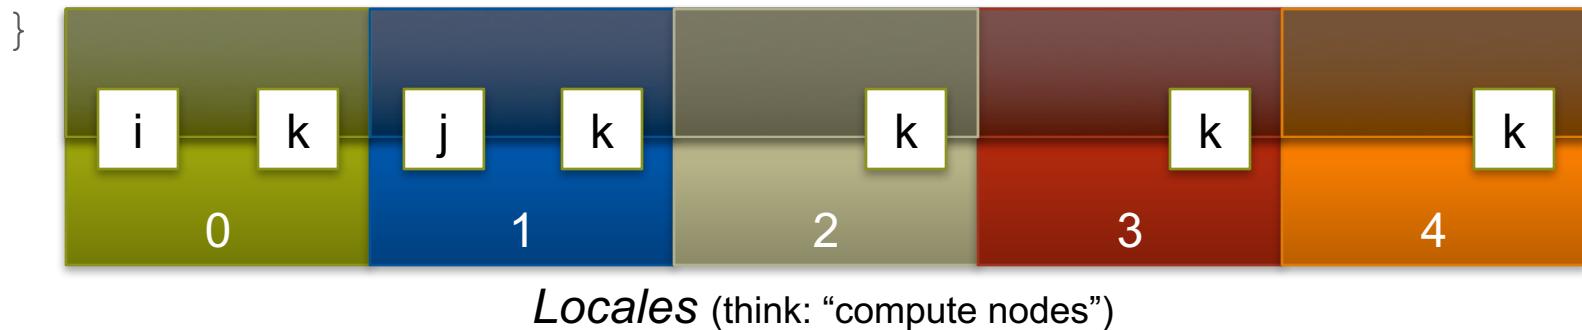
here, *i* and *j* are remote, so  
the compiler + runtime will  
transfer their values

$k = 2*i + j;$

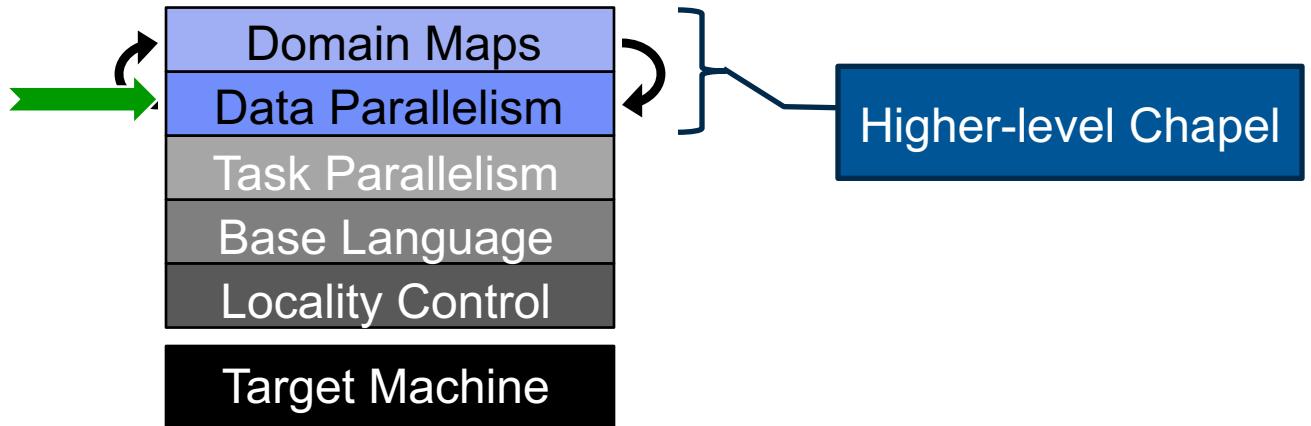


# Chapel: Locality queries

```
var i: int;  
on Locales[1] {  
    var j: int;  
    coforall loc in Locales {  
        on loc {  
            var k: int;  
  
            ...here...           // query the locale on which this task is running  
            ...j.locale...      // query the locale on which j is stored  
        }  
    }  
}  
}  
}
```

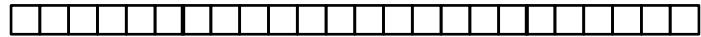


# Data Parallelism

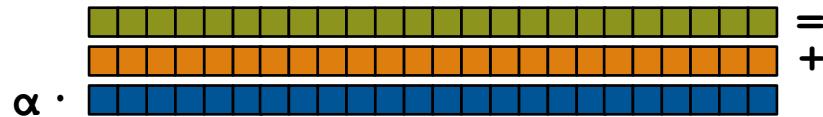


# Data Parallelism By Example: STREAM Triad

```
const ProblemSpace = {1..m};
```



```
var A, B, C: [ProblemSpace] real;
```



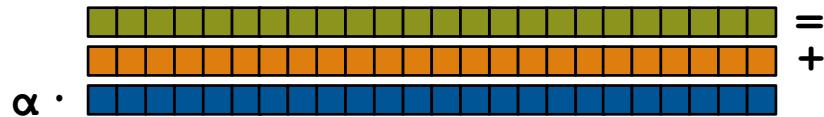
```
forall (a,b,c) in zip(A,B,C) do  
  a = b + alpha*c;
```

# Data Parallelism By Example: STREAM Triad

```
const ProblemSpace = {1..m};
```



```
var A, B, C: [ProblemSpace] real;
```



```
A = B + alpha * C; // equivalent to the zippered forall
```

# Other Data Parallel Features

- **Rich Domain/Array Types:**

- multidimensional
- strided
- sparse
- associative

- **Slicing:** Refer to subarrays using ranges/domains

... `A[2..n-1, lo..#b]` ...  
... `A[ElementsOfInterest]` ...

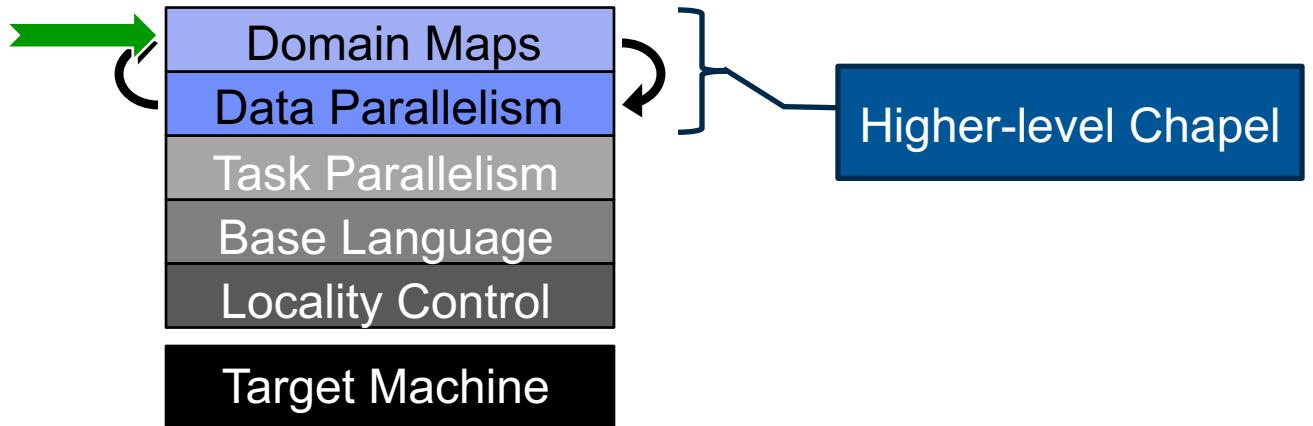
- **Promotion:** Call scalar functions with array arguments

... `pow(A, B)` ... // equivalent to: `forall (a,b) in zip(A,B) do pow(a,b)`

- **Reductions/Scans:** Apply operations across collections

... + **reduce** A ...  
... `myReduceOp reduce A` ...

# Domain Maps

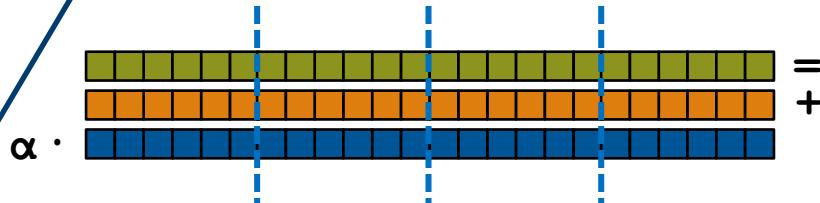


# STREAM Triad: Chapel (multicore)

```
const ProblemSpace = {1..m};
```



```
var A, B, C: [ProblemSpace] real;
```

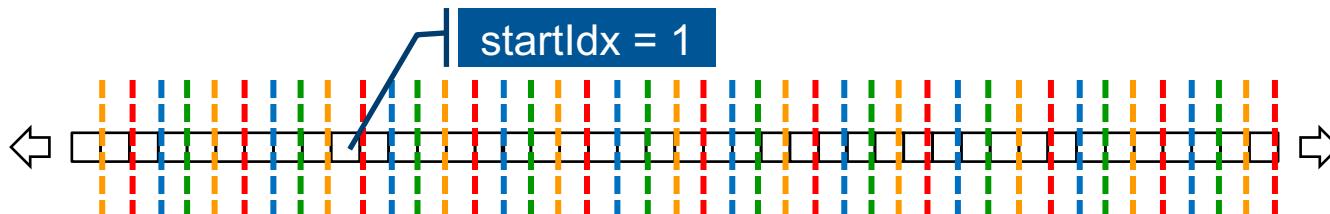


```
A = B + alpha * C;
```

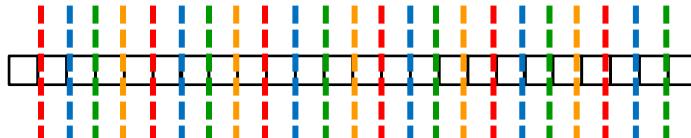
No domain map specified  $\Rightarrow$  use default layout

- current locale owns all domain indices and array values
- computation will execute using local processors only

# STREAM Triad: Chapel (multilocale, cyclic)



```
const ProblemSpace = {1..m}
dmapped Cyclic(startIdx=1);
```

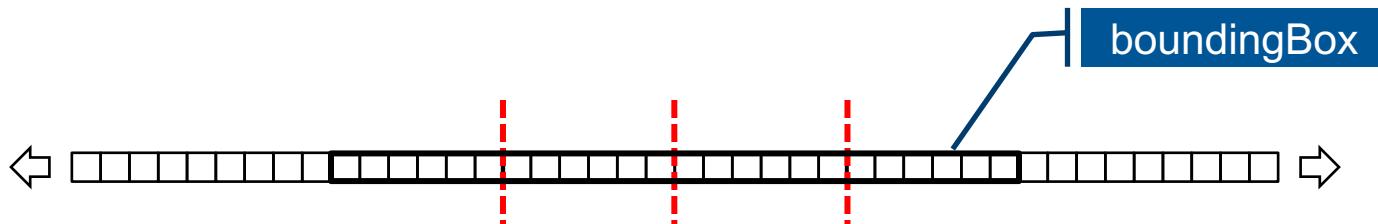


```
var A, B, C: [ProblemSpace] real;
```

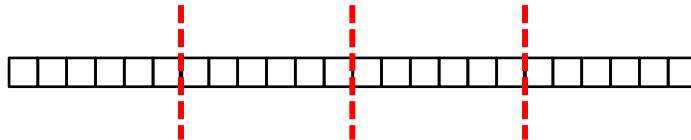


```
A = B + alpha * C;
```

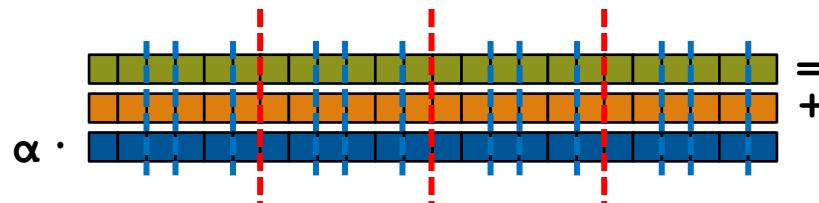
# STREAM Triad: Chapel (multilocale, blocked)



```
const ProblemSpace = {1..m}
dmapped Block (boundingBox={1..m}) ;
```



```
var A, B, C: [ProblemSpace] real;
```



```
A = B + alpha * C;
```

# STREAM Triad: Chapel

MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params,
int myRank, commSize;
int rv, errCount;
MPI_Comm comm = MPI_COMM_WORLD;
MPI_Comm size( comm, &commSize );
MPI_Comm_rank( comm, &myRank );
rv = HPCC_Stream( params, 0 == myRank );
MPI_Reduce( &rv, &errCount, 1, MPI_
return errCount;
}

int HPCC_Stream(HPCC_Params *params,
register int j;
double scalar;
VectorSize = HPCC_LocalVectorSize();
a = HPCC_XMALLOC( double, VectorSize );
b = HPCC_XMALLOC( double, VectorSize );
c = HPCC_XMALLOC( double, VectorSize );

if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d).\n" VectorSize );
        fclose( outFile );
    }
}
```

Chapel

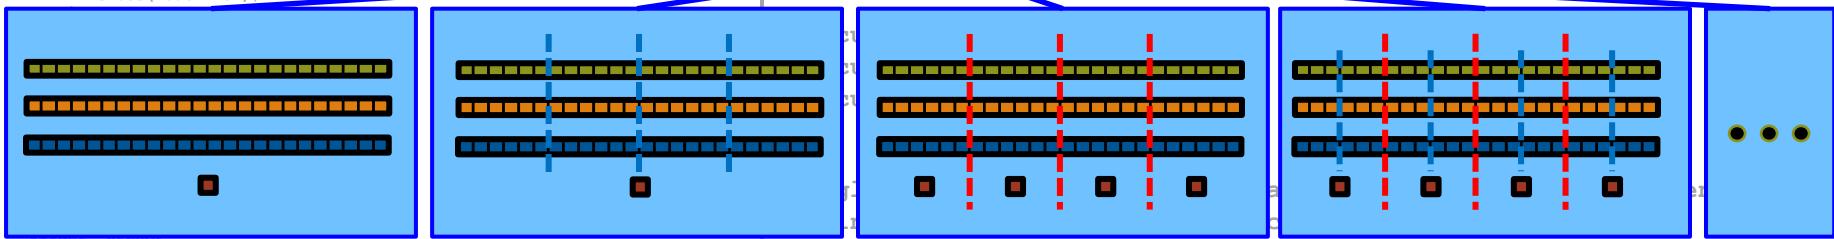
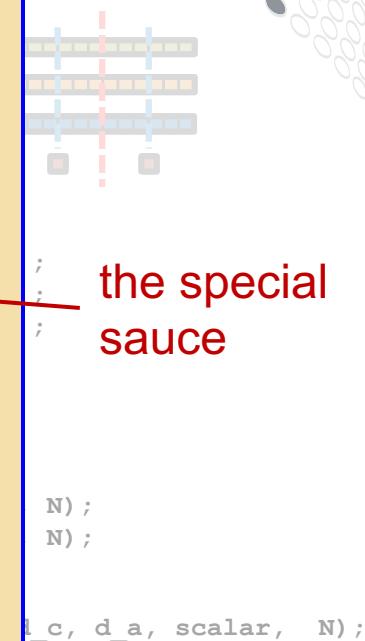
```
config const m = 1000,
alpha = 3.0;

const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

A = B + alpha * C;
```



Philosophy: Good, *top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

# Chapel is Extensible

**Advanced users can create their own...**

- ...domain maps (array layouts and distributions)...
- ...parallel loop schedules...
- ...models of the target architecture...

**...as Chapel code, without modifying the compiler.**

**Why?** To create a future-proof language.

**This has been our main R&D challenge:** How to create a language that does not lock these policies into the implementation without sacrificing performance?



# Language Summary

*Parallel programmers deserve better programming models*

*Higher-level programming models can help insulate algorithms from parallel implementation details*

- yet, without necessarily abdicating control
- Chapel does this via its multiresolution design

*We believe Chapel can greatly improve productivity*

- ...for current and emerging parallel architectures
- ...for HPC users as well as mainstream uses of parallelism

# Outline

- ✓ Motivation for Chapel
- ✓ Survey of Chapel Concepts
- Chapel Project and Characterizations
- Chapel Resources



# Chapel is Portable

- Chapel is designed to be hardware-independent
- The current release requires:
  - a C/C++ compiler
  - a \*NIX environment (Linux, OS X, BSD, Cygwin, ...)
  - POSIX threads
  - RDMA, MPI, or UDP (for distributed memory execution)
- Chapel can run on...
  - ...laptops and workstations
  - ...commodity clusters
  - ...the cloud
  - ...HPC systems from Cray and other vendors
  - ...modern processors like Intel Xeon Phi, GPUs\*, etc.

\* = academic work only; not yet supported in the official release

# Chapel is Open-Source

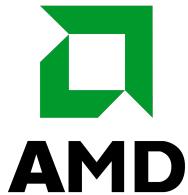
- Chapel's development is hosted at GitHub
  - <https://github.com/chapel-lang>
- Chapel is licensed as Apache v2.0 software
- Instructions for download + install are online
  - see <http://chapel.cray.com/download.html> to get started

# The Chapel Team at Cray (Summer 2016)





# Chapel is a Collaborative Effort



THE GEORGE  
WASHINGTON  
UNIVERSITY  
WASHINGTON, DC



Lawrence Berkeley  
National Laboratory



(and several others...)

<http://chapel.cray.com/collaborations.html>



COMPUTE

STORE

ANALYZE

# Chapel is a Work-in-Progress

- **Currently being picked up by early adopters**
  - over two releases, 3000+ downloads per year
  - Users who try it generally like what they see



# A notable early adopter

## Chapel in the (Cosmological) Wild

1:00 – 2:00

Nikhil Padmanabhan, Yale University Professor, Physics & Astronomy

**Abstract:** This talk aims to present my personal experiences using Chapel in my research. My research interests are in observational cosmology; more specifically, I use large surveys of galaxies to constrain the evolution of the

The image shows a YouTube search results page. At the top left is the YouTube logo. To its right is a search bar with the placeholder text 'Search'. Below the search bar are navigation links: 'Videos', 'Playlists', and 'Channels'. On the left side, there is a thumbnail for a video titled 'CHIUW 2016 keynote: "Chapel in the (Cosmological) Wild", Nikhil Padmanabhan'. The thumbnail shows a man in a white shirt and tan pants standing in front of a projection screen displaying a complex, multi-colored visualization. A progress bar at the bottom of the thumbnail indicates the video is 56:14 minutes long.

**CHIUW 2016 keynote: "Chapel in the (Cosmological) Wild",  
Nikhil Padmanabhan**

Chapel Parallel Programming Language  
1 month ago • 86 views

This is Nikhil Padmanabhan's keynote talk from CHIUW 2016: the 3rd Annual Chapel Implementers and Users workshop. The slides are availabl...

# Chapel is a Work-in-Progress

- **Currently being picked up by early adopters**
  - Last two releases got ~3500 downloads total in a year
  - Users who try it generally like what they see
- **Most current features are functional and working well**
  - some areas need improvements, e.g., error-handling, constructors
- **Performance varies, but is continually improving**
  - shared memory performance is typically competitive with C+OpenMP
  - distributed memory performance tends to be more hit-and-miss
- **We are actively working to address these lacks**

# Outline

- ✓ Motivation for Chapel
- ✓ Survey of Chapel Concepts
- ✓ Chapel Project and Characterizations
- Chapel Resources



# Chapel Websites

## Project page: <http://chapel.cray.com>

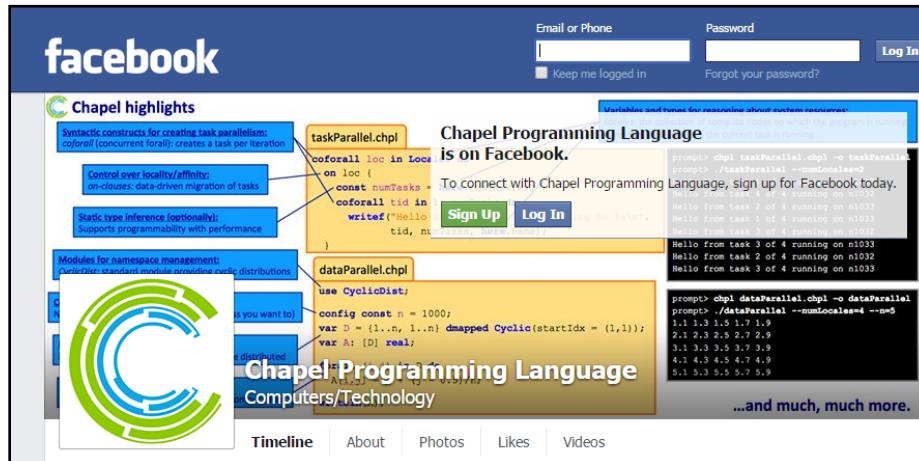
- overview, papers, presentations, language spec, ...

## GitHub: <https://github.com/chapel-lang>

- download Chapel; browse source repository; contribute code

## Facebook: <https://www.facebook.com/ChapelLanguage>

## Twitter: <https://twitter.com/ChapelLanguage>



The screenshot shows the Chapel Programming Language Facebook page. It features a large green and blue 'C' logo. The page highlights include:

- Syntactic constructs for creating task parallelism: `coforall` (concurrent forall), creates a task per iteration.
- Control over locality/affinity: on-edges, data-driven migration of tasks.
- Static type inference (optionally), supports programmability with performance.
- Modules for namespace management: provides a clean way to manage dependencies and distributions.
- Task parallelism: is what you want to do.
- Chapel Programming Language: Computers/Technology

A central callout box contains the text: "Chapel Programming Language is on Facebook. To connect with Chapel Programming Language, sign up for Facebook today." Below this are snippets of Chapel code and a "Sign Up" and "Log In" button.



The screenshot shows the Chapel Language Twitter account (@ChapelLanguage). It has 4 tweets, 1 following, and 19 followers. A recent tweet from March 8, 2016, states: "Chapel is a productive parallel programming language designed for large-scale computing whose development is being led by @cray\_inc". It includes a link to [chapel.sourceforge.net/perf/chapcs/](http://chapel.sourceforge.net/perf/chapcs/). The timeline also shows a chart titled "Binary Trees Shootout Benchmark (n=20)" comparing execution times for different memory allocators.



COMPUTE

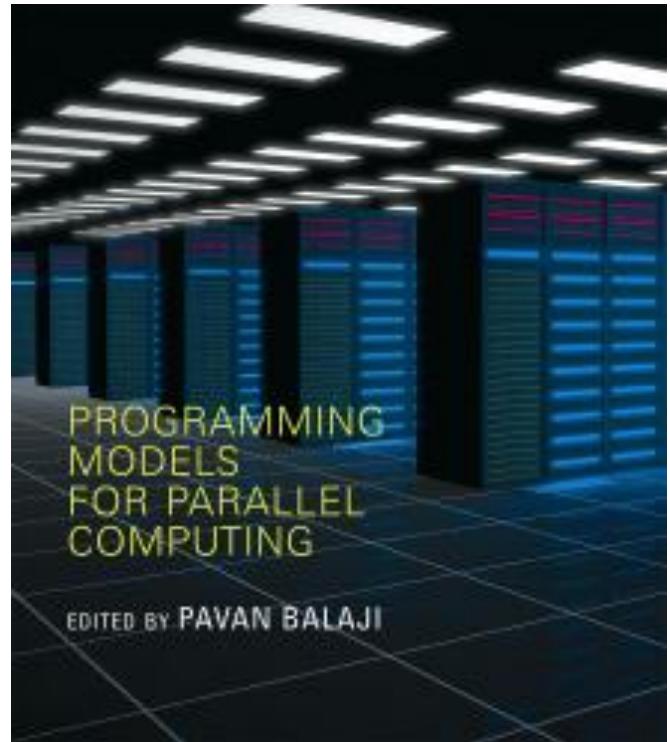
STORE

ANALYZE

# Suggested Reading

Chapel chapter from *Programming Models for Parallel Computing*

- a detailed overview of Chapel's history, motivating themes, features
- published by MIT Press, November 2015
- edited by Pavan Balaji (Argonne)
- chapter is now also available [online](#)



Other Chapel papers/publications available at <http://chapel.cray.com/papers.html>

# Chapel Blog Articles

## [Chapel: Productive Parallel Programming](#), Cray Blog, May 2013.

- *a short-and-sweet introduction to Chapel*

## [Chapel Springs into a Summer of Code](#), Cray Blog, April 2016.

- *coverage of recent events*

## [Six Ways to Say “Hello” in Chapel](#) (parts [1](#), [2](#), [3](#)), Cray Blog, Sep-Oct 2015.

- *a series of articles illustrating the basics of parallelism and locality in Chapel*

## [Why Chapel?](#) (parts [1](#), [2](#), [3](#)), Cray Blog, Jun-Oct 2014.

- *a series of articles answering common questions about why we are pursuing Chapel in spite of the inherent challenges*

## [\[Ten\] Myths About Scalable Programming Languages](#), IEEE TCSC Blog (index available on [chapel.cray.com](#) “blog articles” page), Apr-Nov 2012.

- *a series of technical opinion pieces designed to argue against standard reasons given for not developing high-level parallel languages*

# Chapel Mailing Lists

## low-traffic / read-only:

`chapel-announce@lists.sourceforge.net`: announcements about Chapel

## community lists:

`chapel-users@lists.sourceforge.net`: user-oriented discussion list

`chapel-developers@lists.sourceforge.net`: developer discussions

`chapel-education@lists.sourceforge.net`: educator discussions

`chapel-bugs@lists.sourceforge.net`: public bug forum

(subscribe at SourceForge: <http://sourceforge.net/p/chapel/mailman/>)

## To mail the Cray team:

`chapel_info@cray.com`: contact the team at Cray

`chapel_bugs@cray.com`: for reporting non-public bugs

or use IRC (#chapel on [chat.freenode.net](http://chat.freenode.net)) or StackOverflow



# Current Events: Computer Language Benchmark Game



---

COMPUTE

| STORE

| ANALYZE

# Computer Language Benchmarks Game

**Chapel was recently added to the game:**

**As of Oct 17<sup>th</sup>:**

- **for performance:**

- 1 top entries: pidigits
- 2 top-5 entries: meteor, thread-ring
- 2 top-10 entries: fannkuch-redux, chameneos-redux
- 3 top-20 entries: n-body, spectral-norm, binary-trees

- **for code compactness:**

- 2 top entries: n-body, thread-ring
- 2 top-5 entries: spectral-norm, pidigits
- 4 top-20 entries: mandelbrot, regex-dna, chameneos-redux, meteor

## The Computer Language Benchmarks Game

64-bit quad core data set

Will your toy benchmark program be faster if you write it in a different programming language? It depends how you write it!

### Which programs are fast?

Which are succinct? Which are efficient?

<u>Ada</u>	<u>C</u>	<u>Chapel</u>	<u>Clojure</u>	<u>C#</u>	<u>C++</u>
<u>Dart</u>	<u>Erlang</u>	<u>F#</u>	<u>Fortran</u>	<u>Go</u>	<u>Hack</u>
<u>Haskell</u>	<u>Java</u>	<u>JavaScript</u>	<u>Lisp</u>	<u>Lua</u>	
<u>OCaml</u>	<u>Pascal</u>	<u>Perl</u>	<u>PHP</u>	<u>Python</u>	
<u>Racket</u>	<u>Ruby</u>	<u>JRuby</u>	<u>Rust</u>	<u>Scala</u>	
	<u>Smalltalk</u>	<u>Swift</u>	<u>TypeScript</u>		



# Computer Language Benchmarks Game

Chapel was recently added to the game:

As of Oct 17th:

## The Computer Language Benchmarks Game

64-bit quad core data set

- for Chapel
- We want **easy answers**, but easy answers are often incomplete or wrong. You and I know, there's more we should understand:

stories      details      fast?      conclusions

{ for researchers }

4 top-20 entries: mandelbrot, regex-dna, chameneos-redux, meteor

Racket	Ruby	JRuby	Rust	Scala
Smalltalk	Swift	TypeScript		



COMPUTE

STORE

ANALYZE



# Chapel: Productive Parallel Programming at Scale

Questions?



---

COMPUTE

| STORE

| ANALYZE